Copyleaks

Analysis Report

Scan details

Scan Time	Total Pages	Total Words
May 7th, 2026 at 12:52 UTC	10	2420

Plagiarism Detection and AI Detection Report

Research ppr_2026 (10).pdf

Plagiarism Detection

0%

Plagiarism types	Text coverage	Words
Identical	0%	0
Minor Changes	0%	0
Paraphrased	0%	0
Excluded		
Omitted Words		1

AI Detection

0%

	Text coverage	Words
AI Text	0%	0
Human Text	100%	1
Excluded		
Omitted Words		1

Plagiarism

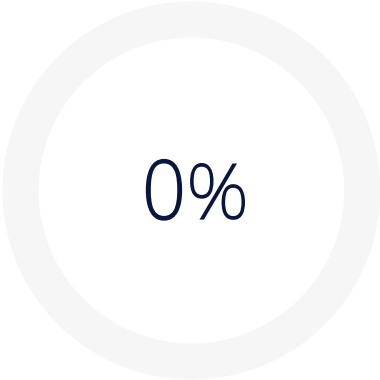

Results (0)

*Results may not appear because the feature has been disabled.

Private Cloud Hub	Shared Data Hub	Filtered / Excluded
0	0	0
Internet Sources	**AI Source Match**	**Current Batch**
0	0	0

Plagiarism types	Text coverage	Words
Identical	0%	0
Minor Changes	0%	0
Paraphrased	0%	0
Excluded		
Omitted Words		1

About Plagiarism Detection

Our AI-powered plagiarism scans offer three layers of text similarity detection: Identical, Minor Changes, and Paraphrased. Based on your scan settings we also provide insight on how much of the text you are not scanning for plagiarism (Omitted words).

Identical

One to one exact word matches. Learn more

Minor Changes

Words that hold nearly the same meaning but have a change to their form (e.g. “large” becomes “largely”). Learn more

Paraphrased

Different words that hold the same meaning that replace the original content (e.g. 'large' becomes 'big') Learn more

Omitted Words

The portion of text that is not being scanned for plagiarism based on the scan settings. (e.g. the 'Ignore quotations' setting is enabled and the document is 20% quotations making the omitted words percentage 20%) Learn more

Copyleaks Shared Data Hub

Our Shared Data Hub is a collection of millions of user-submitted documents that you can utilize as a scan resource and choose whether or not you would like to submit the file you are scanning into the Shared Data Hub. Learn more

Filtered and Excluded Results

The report will generate a complete list of results. There is always the option to exclude specific results that are not relevant. Note, by unchecking certain results, the similarity percentage may change. Learn more

Current Batch Results

These are the results displayed from the collection, or batch, of files uploaded for a scan at the same time. Learn more

Transformer-Based Natural Language Processing

Vibha
Department of Computer Science
Chitkara University
Rajpura, Punjab, India
{vibha2521.be22@chitkara.edu.in }

Ridhima
Department of Computer Science
Chitkara University
Rajpura, Punjab, India
{ridhima2151.be22@chitkara.edu.in

ABSTRACT

NLP has had various phases throughout history, from rule-based techniques to a modern approach of deep learning models. Work based on rules has depended on predefined standard rules which constrained it in language complexity. Unlike the rule-based approach today, NLP models are developed utilizing machine and deep learning techniques that can easily interpret human languages. Transformer based models have been amongst the most game-changing ones in the last years due to their unmatched context capturing capabilities.

In this paper, it discusses a few famous models used in NLP tasks such as BERT and GPT transformers, and gives you a short background of NLP. It also covers downstream NLP models like sentiment analysis, machine translation, and question-answering. Then they perform an experiment over IMDB movie reviews using Logistic Regression and Random Forest in the context of sentiment analysis.

The experiment results show that traditional machine learning models can provide strong baselines.However, transformers seem better suited for learning from context clues. However, there are certain considerations that need to be taken into account in relation to computational power, biases, and explainability in the future.

Keywords— Natural Language Processing, Transformers, BERT, GPT, Deep Learning, Sentiment Analysis, Machine Learning, Text

I. INTRODUCTION

Natural Language Processing (NLP) is a field that puts artificial intelligence to work in the interpretation and analysis of human languages. It is a study which includes computer science, linguistics and artificial intelligence. At first what was done in NLP was the development by experts of rule based systems which did very well in very specific situations. Later we saw in to play a role statistical methods and machine learning which improved what computers could do in the area of NLP by getting them to identify patterns in data, though in practice they did not do a great job at understanding the full context and meaning. We have seen great progress in NLP with the introduction of deep learning algorithms. We had Recurrent Neural Networks (RNNs) and Long Short Term Memory (LSTM) networks which did a better job at what they term sequence processing but had issues with which we still struggle with long texts. The issue was solved with the use of self-attention in Transformer models which in turn enabled the model to pay attention to important words in a sentence. Also as a result transformers have become the base of many NLP applications. At present the field of NLP is used in virtual assistants, translation, social networks' sentiment analysis, education, medicine, customer service and so on. Also due to the fact that language is a key element of communication, constant improvement in NLP is a must for real world applications. This in turn has put forth great research interest in transformer based approaches.

II. BACKGROUND

In general, the development of Natural Language Processing can be divided into three main stages. The first stage involved the use of rule-based techniques in Natural Language Processing, during which the analysis of languages was performed on the basis of grammatical rules and handcrafted dictionaries [13]. While being a simple solution, it was not flexible enough to cope with the variety of languages encountered in practice.

The second stage implied the application of statistical approaches to NLP. In this period, n-gram techniques and probabilistic models were used [4]. Though learning from data was possible in this approach, it heavily relied on manually crafted features, failing to capture the context.

The third stage witnessed the advent of Deep Learning techniques that made language understanding much more accurate due to the ability to model complex dependencies within the text. However, models based on the RNN [6] and LSTM architecture [7] were prone to the vanishing gradient problem that limited their capability to capture long-range dependencies. This issue was successfully addressed using attention mechanisms, presented in [2].

III. METHODOLOGY

This study we present a mixed method which includes a literature review and relative analysis. We looked at what we term colorful exploration papers, reports which covered BERT[1] and GPT[3] in the fields of NLP development, motor infrastructure, and their functions. We looked at these through the lens of traditional rule based approaches, machine learning methods, deep learning techniques, and motor based models. Also we did a relative study which looked at issues like accuracy, computational efficiency, scalability, and context which these models do or do not have. In terms of rule based approaches we looked at the use of verbal rules. For machine learning we looked at the use of classifiers and feature extraction. As for motor based models we looked at issues of tuneable parameters, pre training, fine tuning. Also we looked at the pretraining and fine tuning processes[2]. For the experiments, we used the IMDB movie review dataset to classify sentiment. After converting the text data into numerical form using TF-IDF, we compared the performance of Logistic Regression and Random Forest algorithms. This approach allowed us to combine theoretical analysis with practical testing for a comprehensive understanding

IV. TRANSFORMER ARCHITECTURE

Transformers are neural networks with a structure that has been specifically designed to handle natural language, matching or even exceeding the performance of traditional sequence-based models. A core component of transformer networks is so-called self-attention, where each word in a sentence is weighted by the network in terms of importance to the overall sentence.

Most transformer models are built with two parts: an encoder and a decoder. The encoder processes the input sequence, and the decoder then generates the output sequence. Because the transformer architecture does not force the model to keep the order of input words, the model needs to be provided with positional information. This is achieved by adding positional encoding. The transformer and similar models have achieved state-of-the-art results in many NLP tasks [12].

Recent transformer-based models like BERT [1] and GPT [3] have garnered remarkable attention in the NLP community. Commonly, these deep models are first pre-trained on a large corpus of text and then further fine-tuned on a task-specific dataset for adaptation to new downstream applications.

V. TRANSFORMER MODELS AND VARIANTS

BERT seems to be the front runner of all the transformer models released recently. Being able to see input text from both sides helps BERT understand context better which is why it's especially useful for tasks like sentiment analysis and question answering. On the other hand GPT

uses auto-regression to fill in text by trying to predict the next word based on what has already been generated, hence it excels in text generation and in conversation applications. It's also worth noting that GPT achieves improved performance through large scale training. In addition to BERT and GPT, a slew of additional variants have been released that either perform better or are more scalable. Models like RoBERTa, T5, ALBERT, and DistilBERT have all been released with improved performance and scalability, further demonstrating that transformer research is not only about models getting bigger but also about models becoming more useful and logical

VI. Applications of NLP

Natural Language Processing is being increasingly used in a variety of real-world domains. Machine translation of text from one language to another is a particularly common use of NLP. Recent translation systems based on transformer architectures [2] have outperformed prior approaches in terms of capturing context.

Sentiment analysis is another common application where you want to know if some text is positive, negative or neutral. This is used for product reviews, social media monitoring and customer feedback analysis. The models based on the transformer architecture improve such systems by better understanding the context.

NLP is also used in other applications including chatbots, virtual assistants, speech recognition, text summarization and question answering systems. Note that with the growing need for human machine interaction, NLP has become a fundamental part of current information technology systems.

VII. Experimental Evaluation

A. Experimental Setup

Sentiment classification is a problem where you try to determine if someone has expressed something positive or negative in given data. In this work, I tried to see how good a few common machine learning techniques are for a simple sentiment classification task where data consists only of text and the task is to classify the data between positive and negative. I used part of the well-known dataset IMDB movie reviews in this work. The dataset is widely used in related research

B. Data Analysis

Reviews were roughly balanced, containing almost equal numbers of positive and negative reviews. During the exploratory data analysis positive reviews often included words as "excellent," "great," or "amazing," while negative reviews used words such as "bad," "boring" or "poor." There are patterns in this textual data that can be learned by machine learning algorithms [4] .

C. Model Evaluation

In this paper, two machine learning models, namely, Logistic Regression and Random Forest are used. Text was preprocessed and transformed into numerical feature vectors using TF-IDF. Such techniques are generally employed in NLP tasks for classification purposes [12]. Accuracy obtained by the Logistic Regression classifier was 87%, and that of Random Forest was 84.92%, thus indicating that the performance of the Logistic Regression model is better for the current dataset.

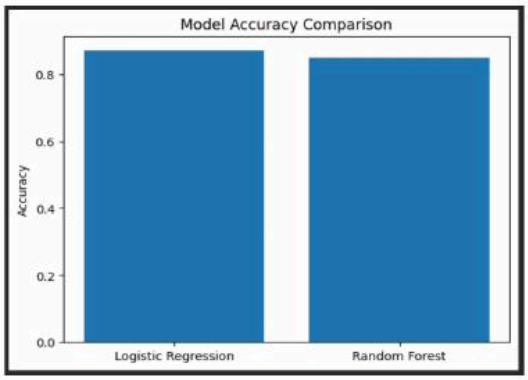

Fig1.Model Accuracy comparison

D. Confusion Matrix Analysis

According to the so-called confusion matrix the model correctly identified 4271 negative reviews and 4429 positive reviews. 690 negative reviews and 610 positive reviews were misclassified, however

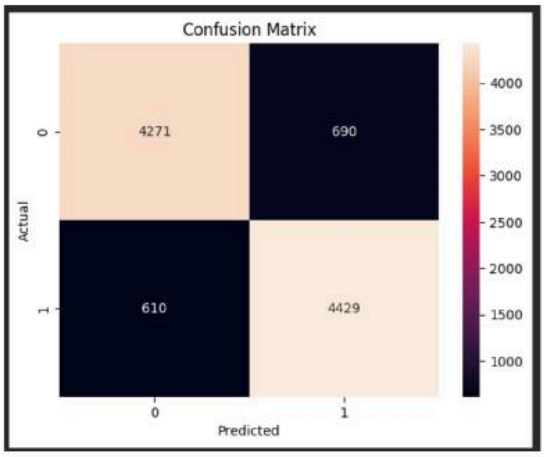

Fig 2.Confusion Matrix Analysis

E. Visualization

To make the results more understandable, accuracy comparison chart and confusion matrix heatmap were presented to viewers..

F. Applications

The findings have implications for using Natural Language Processing (NLP) approaches for analysis of reviews of products and services, customer feedback, and social media monitoring.

G. Key Findings

In the experiment NLP methods were able to achieve decent results in the sentiment classification task. Simple machine learning models like Logistic Regression even slightly outperformed a Random Forest, serving as a strong baseline.

H. Conclusion of Experiment

State-of-the-art results for sentiment analysis have been reported by traditional machine learning approaches, yet these results can be still outperformed by transformer-based models, such as BERT [1] and GPT [3], that utilize powerful contextual models.

VIII. DISCUSSION

Results show that current state-of-the-art models, based on transformer architecture, understand context and semantic relationships better than traditional models. However, they require significantly more resources to run.

While improving accuracy is a natural goal for neural networks, current research also strives to reduce computational cost, decrease model size and increase interpretability.

Bias. It's a problem that we have noticed when working with the huge data sets that were used for the pretraining of the used models. Although this phenomenon has already been observed by other researchers for similar large language models like GPT [15], it doesn't make it any less problematic.

IX. Challenges and Limitations

Transformer models have been revolutionary in terms of achieving state-of-the-art results on many benchmarks. However, one of the major shortcomings of the transformers is their expensive computation. Training and running transformer models requires resources that are out of reach for many small companies.

Another disadvantage is that the models are poorly interpretable, although they still reach good results [12]. However, there are many applications in health care and even in the legal system, where interpretability is of importance.

Another challenge is that of language ambiguity in preannotated training data.Natural human languages are unclear and transformers many time find it difficult to understand them [13].

X. Future Directions

Efficiency, Interpretability, and Adaptability are likely to be the key elements of the future research in NLP. Some efficient models like DistilBERT [11] and ALBERT [10] have already been used with high quality and less computational power.

The other is called few-shot learning which is a task carried out with a few labeled samples [15] and is gaining more interest these days. This will be particularly helpful in the case of under-resourced languages. One of the important issues is the multimodality which involves the interaction between text and image, or even sound and text.

Factors of fairness and accountability are always important. NLP systems should be objective, transparent and trustworthy in the future.

XI. Conclusion

NLP has gone from a set of rules to a model of transformer, making a lot of progress in understanding human languages and natural languages [13]. Examples of such models include BERT [1] and GPT [3] that have made significant contributions to many applications, such as machine translations, text classifications or language generation.

The results presented in this paper demonstrate the effectiveness of conventional ML methods to sentiment analysis. However, transformer models offer a greater understanding in context [2]. However, there are important questions that remain to be answered. The challenges in this regard are computational efficiency, model bias [14] and model interpretability.

To conclude, transformer-based NLP is a phenomenal advancement in the field of AI. It's clear that it will continue to improve over the years.

XI. REFERENCES

[1] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” arXiv preprint arXiv:1810.04805, 2018.

[2] Ashish Vaswani et al.,“Attention Is All You Need,”in Advances in Neural Information Processing Systems (NeurIPS), vol. 30, 2017.

[3] Alec Radford, K. Narasimhan, T. Salimans, and I. Sutskever,“Improving Language Understanding by Generative Pre-Training,”OpenAI, 2018.

[4] Tomas Mikolov, K. Chen, G. Corrado, and J. Dean,“Efficient Estimation of Word Representations in Vector Space,”arXiv preprint arXiv:1301.3781, 2013.

[5] Andrew Maas et al.,“Learning Word Vectors for Sentiment Analysis,”in Proceedings of the 49th Annual Meeting of the ACL, 2011.

[6] Sepp Hochreiter and Jürgen Schmidhuber,“Long Short-Term Memory,”Neural Computation, vol. 9, no. 8, pp. 1735–1780, 1997.

[7] Kyunghyun Cho et al.,“Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation,”in EMNLP, 2014.

[8] Yinhan Liu et al.,“RoBERTa: A Robustly Optimized BERT Pretraining Approach,”arXiv preprint arXiv:1907.11692, 2019.

[9] Colin Raffel et al.,“Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer,”Journal of Machine Learning Research, vol. 21, pp. 1–67, 2020.

[10] Zhenzhong Lan et al.,“ALBERT: A Lite BERT for Self-supervised Learning of Language Representations,”ICLR, 2020.

[11] Victor Sanh et al.,“DistilBERT: A Distilled Version of BERT,”arXiv preprint arXiv:1910.01108, 2019.

[12] Yoav Goldberg,“A Primer on Neural Network Models for Natural Language Processing,”Journal of Artificial Intelligence Research, vol. 57, pp. 345–420, 2016.

[13] Christopher Manning and Hinrich Schütze,Foundations of Statistical Natural Language Processing, MIT Press, 1999.

[14] Ian Goodfellow, Y. Bengio, and A. Courville,Deep Learning, MIT Press, 2016.

[15] Tom Brown et al.,“Language Models are Few-Shot Learners,”in Advances in Neural Information Processing Systems (NeurIPS), vol. 33, 2020.

AI Content

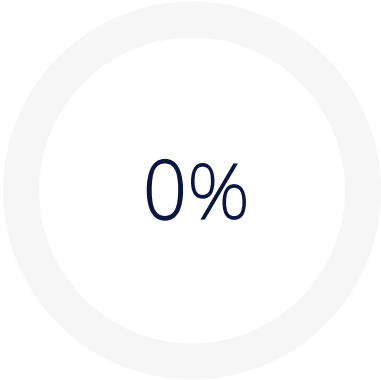

	Text coverage	Words
AI Text	0%	0
Human Text	100%	1
Excluded		
Omitted Words		1

About AI Detection

Our AI Detector is the only enterprise-level solution that can verify if the content was written by a human or generated by AI, including source code and text that has been plagiarized or modified. Learn more

AI Text

A body of text that has been generated or altered by AI technology. Learn more

Human Text

Any text that has been fully written by a human and has not been altered or generated by AI. Learn more

Copyleaks AI Detector Effectiveness

Credible data at scale, coupled with machine learning and widespread adoption, allows us to continually refine and improve our ability to understand complex text patterns, resulting in over 99% accuracy–far higher than any other AI detector–and improving daily. Learn more

Ideal Text Length

The higher the character count, the easier for our technology to determine irregular patterns, which results in a higher confidence rating for AI detection. Learn more

Reasons It Might Be AI When You Think It's Not

The AI Detector can detect a variety of AI-generated text, including tools that use AI technology to paraphrase content, auto-complete sentences, and more. Learn more

User AI Alert History

Historical data of how many times a user has been flagged for potentially having AI text within their content. Learn more

AI Logic

The number of times a phrase was found more frequently in AI vs human text is shown according to low, medium, and high frequency. Learn more